

Brief Intro to Machine Learning

What is Machine Learning?

- From [Simon 83]
 - Learning denotes **changes** in the system that are adaptive in the sense that they enable the system to do the same task or tasks drawn from the same population **more effectively** the next time
- From [Nilsson 96]
 - A machine learns whenever it **changes** its structure, program, or data in such a manner that its expected future performance **improves**


Examples

- Character recognition
- raw data: image
- class: numerals, English (Chinese, etc.) characters




Examples

- **Speech recognition**
- raw data: speech signal
- class: spoken words



Examples

- Face detection
- raw data: image
- class: face/non-face



Examples

- **Fingerprint identification**
 - raw data: fingerprint image
 - class: known/unknown person



Examples

- Document classification
- raw data: (web) document
- class: semantic categories

User Guide

The User Guide covers all of pandas by topic area. Each of the subsections introduces a topic (such as “working with missing data”), and discusses how pandas approaches the problem, with many examples throughout.


Users brand-new to pandas should start with [10 minutes to pandas](#).

Further information on any specific method can be obtained in the [API reference](#).

- IO tools (text, CSV, HDF5, ...)
 - CSV & text files
 - JSON
 - HTML
 - Excel files
 - OpenDocument Spreadsheets
 - Clipboard

Why can this be solved?

- Is this an apple?




Training examples

Apples



Not Apples




Some terminologies

- **Model:** the collection of parameters you are trying to fit
- **Data:** what you are using to fit the model
- **Target:** the value you are trying to predict with your model
- **Features:** attributes of your data that will be used in prediction
- **Methods:** algorithms that will use your data to fit a model

Features

Automatic fish sorting



Different types of fish differ in length, lightness, width, number and shape of fins, position of the mouth, etc

Also there are variations in lighting, position of the fish on the conveyer, etc

- noise
- needs pre-processing


Example of Pre-processing

- Remove background



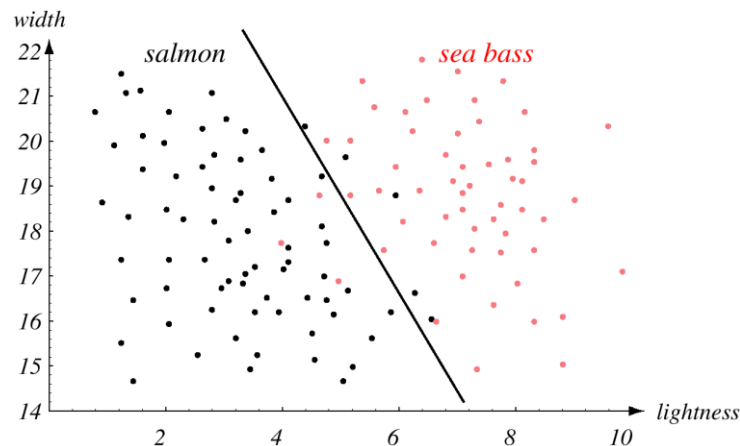
Example of Pre-processing

- Adjust the light level



Feature extraction

Consider each fish as a **point** in some **feature space**



2-dimensional feature space: $x = [x_1, x_2]^T$


Make sure you have a **good** set of features to work with!

Example (bad feature)


Location of a fish on the conveyer belt

Examples of features

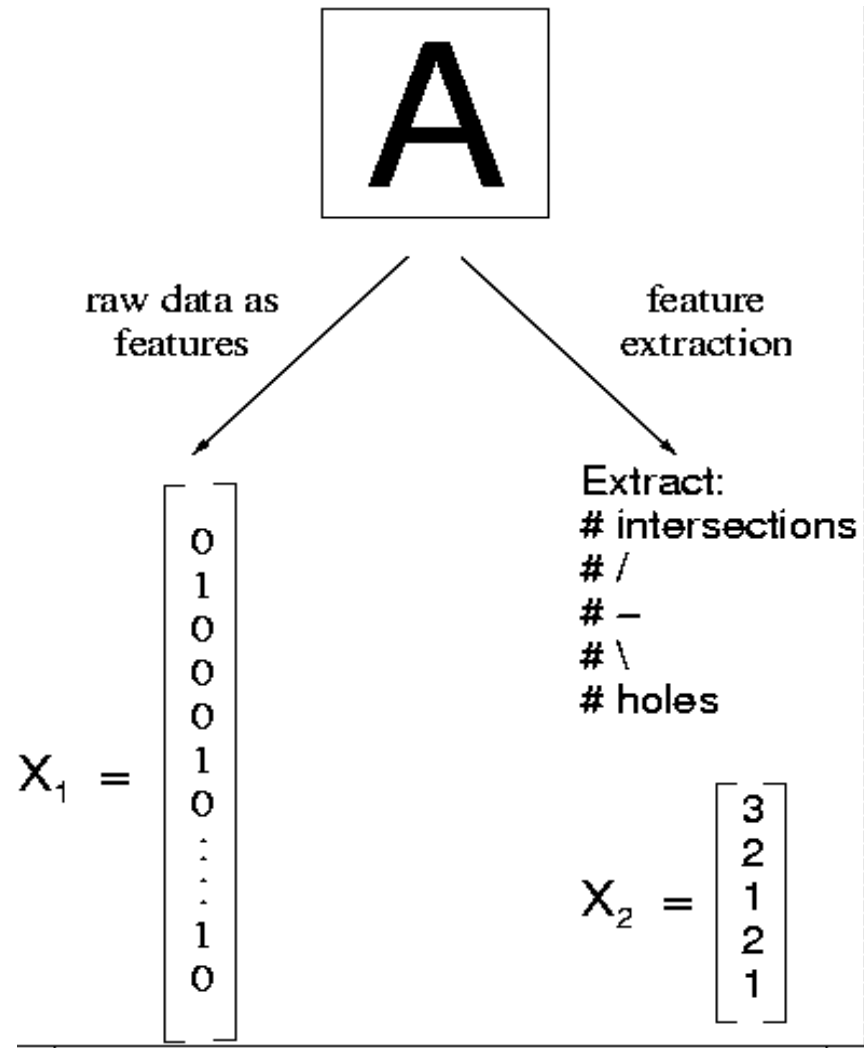
waveform



character




Examples of features



Learning

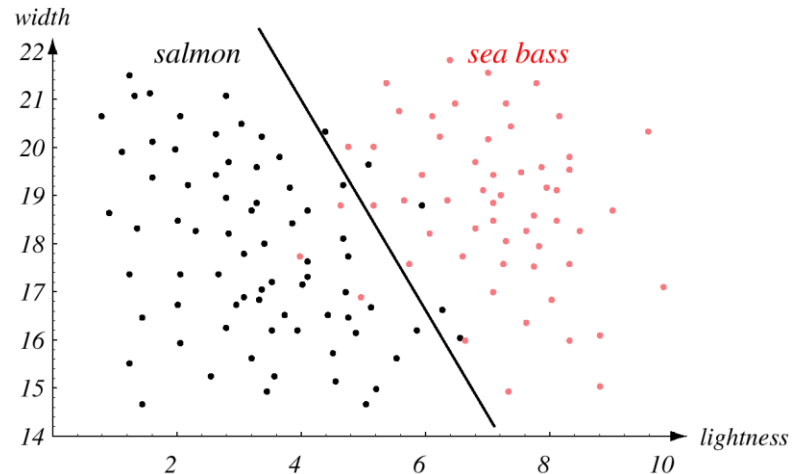
Training sample

- tell me which species a particular fish belongs to
(data collection)



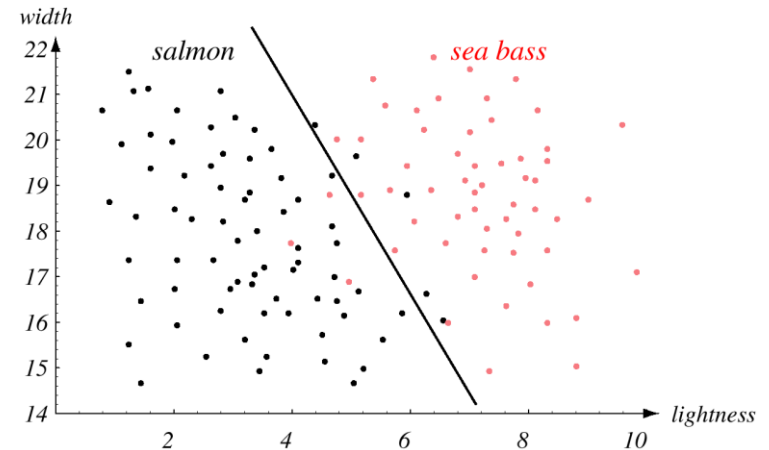
Partition the feature space into 2 regions, one for each type of fish

- decision boundary

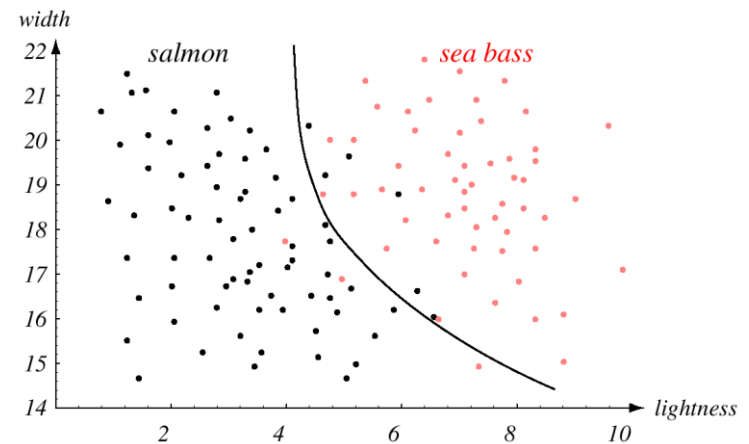


Classifier

Linear classifier




Can also use a more complicated
decision boundary
e.g. **quadratic** classifier




Classifier


can be even more complicated



**linear
boundaries**




**quadratic
boundaries**



**more complicated
boundaries**

How to handle new images?

- Use the trained classifier to **classify** the new image





Question: How to measure the classifier performance?

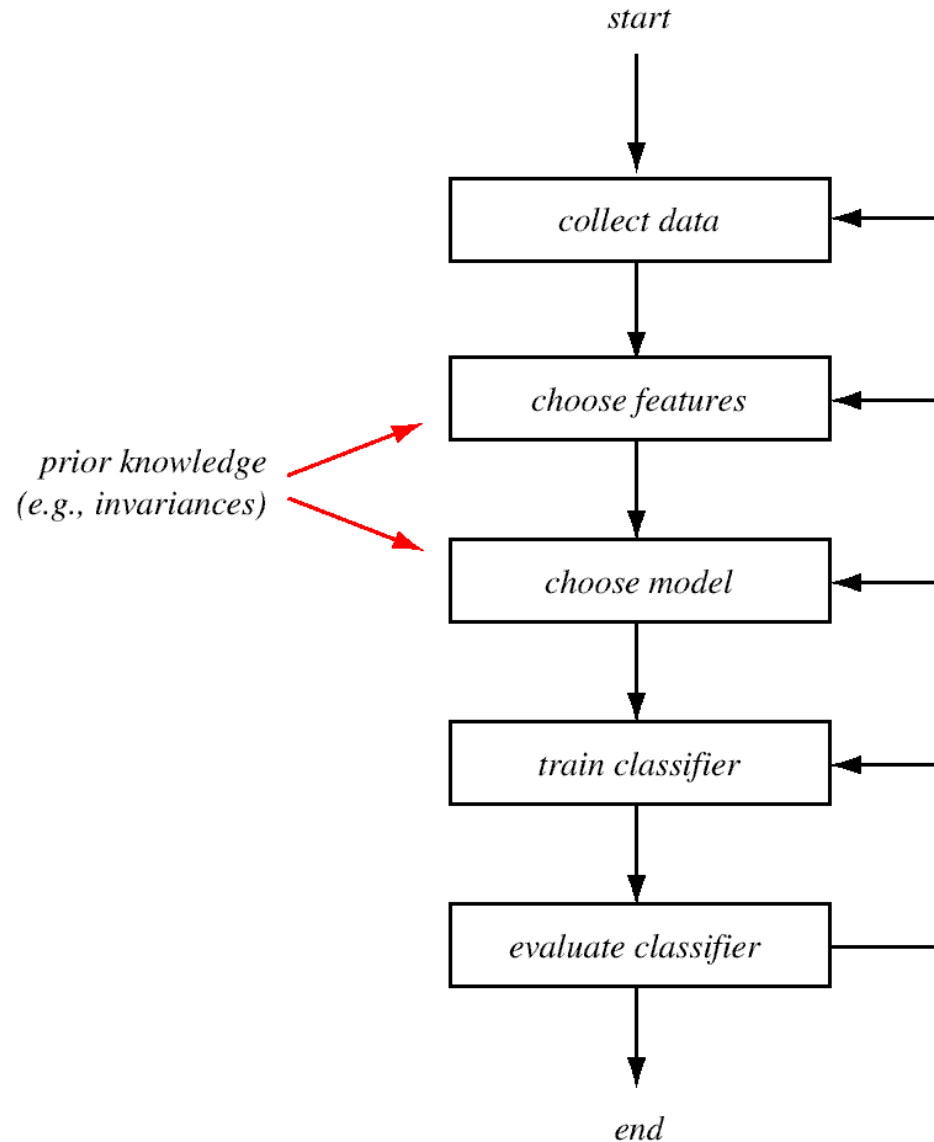
- classification error rate
% patterns that are assigned to the wrong category
- other aspects may be important too
e.g. computational complexity
e.g. user-friendliness

How to handle new images?

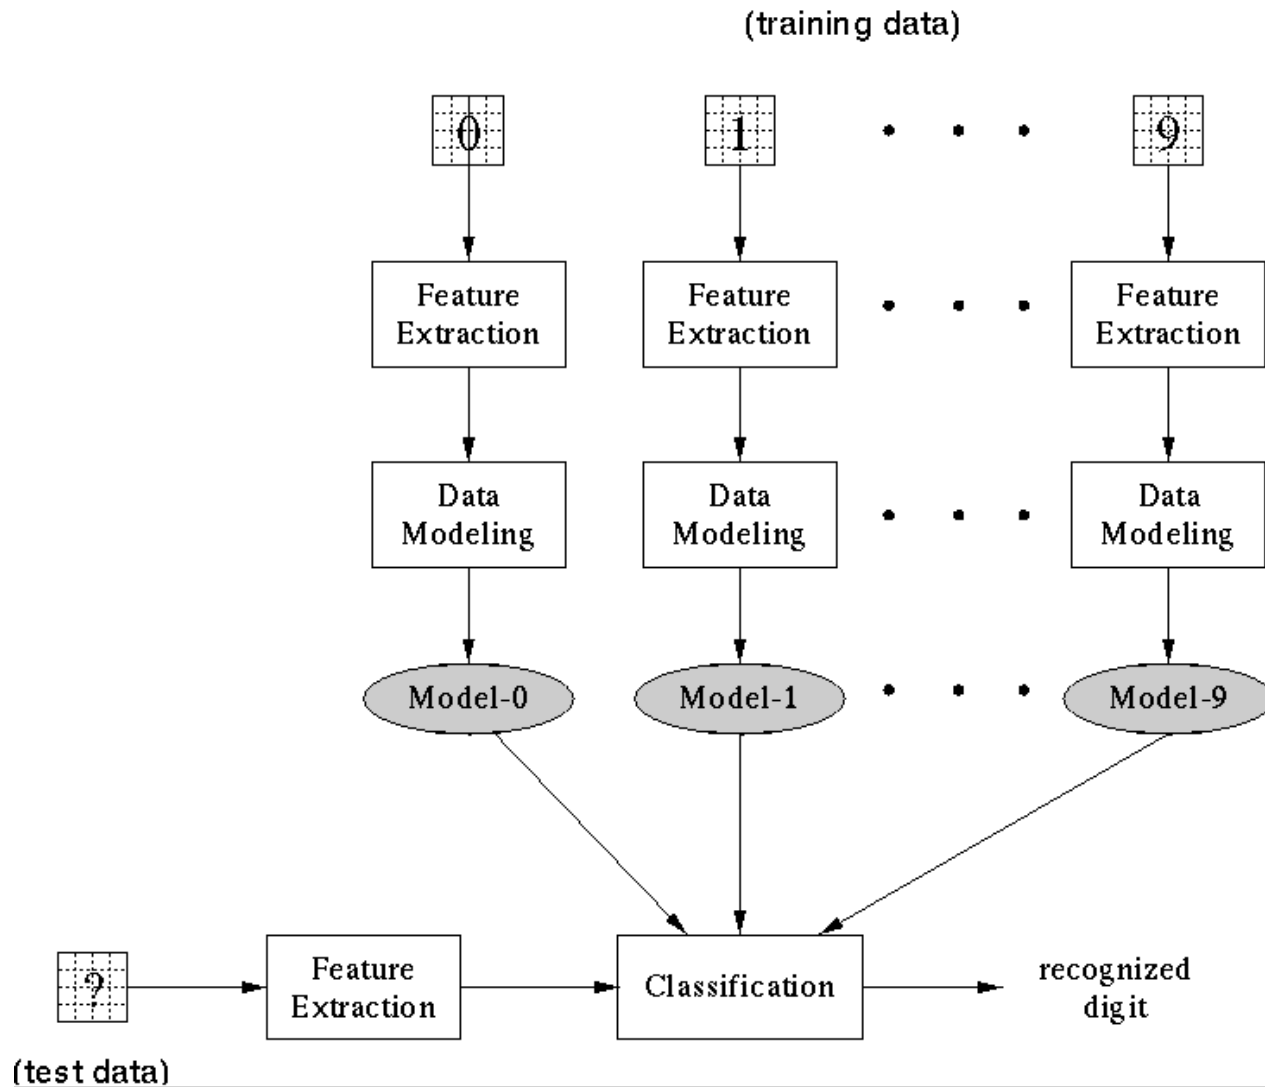
- Question: will the classifier work for this unseen fish?
 - Issue of generalization
- Example:
 - Which model is better?



Classifier design




Example: OCR



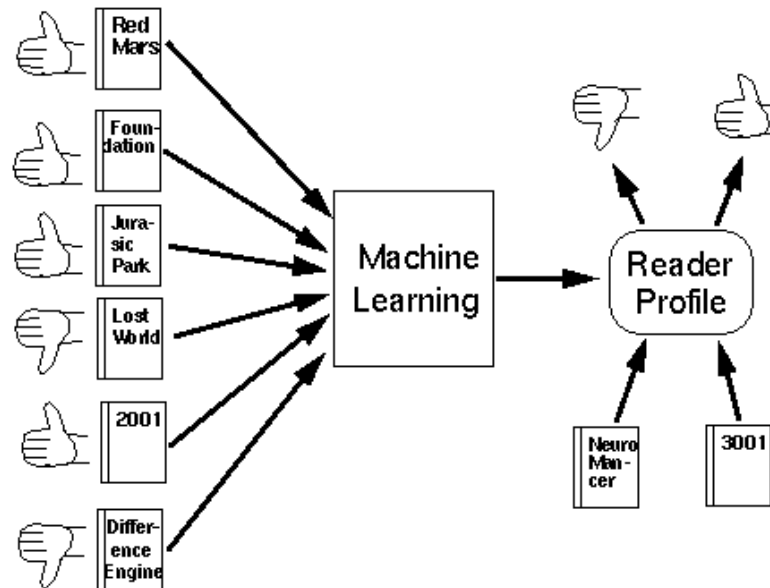

Learning paradigm: supervised learning

- The learner is provided with a set of **inputs** together with the corresponding desired **outputs**




- Has a **teacher**
- Example:
 - teaching kids to recognize different animals
 - graded examinations with correct answers provided

Example: Recommender System



Learning paradigm: unsupervised learning

- Training examples as **input** patterns, with no associated output



- **No teacher**
- similarity measure exists to detect groupings / clusterings


Example: Novelty Detection

- Network intrusion detection

File Logs Settings Help								
Application	Protocol	Local Address	Remote Address	State	Creation Time	Rx [Bytes]	Type	Tx [Bytes]
SYCHOST.EXE	TCP	all:135	-----	Listening	27/Feb/2007 09:41:27	0	0	0
SYCHOST.EXE	TCP	all:3389	-----	Listening	27/Feb/2007 09:41:33	0	0	0
SYCHOST.EXE	TCP	all:1025	-----	Listening	27/Feb/2007 09:41:29	0	0	0
SYSTEM	TCP	192.168.123.128:139	-----	Listening	27/Feb/2007 09:41:30	0	0	0
PERSFW.EXE	TCP	all:44334	-----	Listening	27/Feb/2007 09:41:43	0	0	0
PERSFW.EXE	TCP	all:44334	localhost:1986	Connected In	27/Feb/2007 16:45:42	2931	<attack>	152516
SYCHOST.EXE	TCP	all:5000	-----	Listening	27/Feb/2007 09:42:15	0	0	0
LSASS.EXE	TCP	all:27155	-----	Listening	27/Feb/2007 09:42:17	0	0	0
PUTTY.EXE	TCP	all:1898	uststu1.ust.hk:22	Connected Out	27/Feb/2007 15:54:54	445204	<normal>	141398
CCAPP.EXE	TCP	localhost:1035	-----	Listening	27/Feb/2007 09:42:23	0	0	0
SSHCLIENT.EXE	TCP	all:1639	lcpu2.cse.ust.hk:22	Connected Out	27/Feb/2007 11:43:48	135198	<normal>	3764
ICQ.EXE	TCP	all:21470	-----	Listening	27/Feb/2007 14:41:02	0	0	0
ICQ.EXE	TCP	all:1759	64.12.24.205:5190	Connected Out	27/Feb/2007 14:41:04	59285	<normal>	12191
MSNMSG.R.EXE	TCP	all:1722	by2msg1104003.ph...	Connected Out	27/Feb/2007 14:38:33	176880	<normal>	129012
PFWADMIN.EXE	TCP	all:1986	localhost:44334	Connected Out	27/Feb/2007 16:45:42	278215	<normal>	2931
LSASS.EXE	UDP	all:500	-----	Listening	27/Feb/2007 09:41:40	0	0	0
SYSTEM	UDP	192.168.123.128:138	-----	Listening	27/Feb/2007 09:41:30	205	0	40995
LSASS.EXE	UDP	localhost:1033	-----	Listening	27/Feb/2007 09:42:23	0	0	2
SYSTEM	UDP	192.168.123.128:137	-----	Listening	27/Feb/2007 09:41:30	1088	0	17982
PERSFW.EXE	UDP	all:44334	-----	Listening	27/Feb/2007 09:41:43	0	0	0
IEXPLORE.EXE	UDP	localhost:1600	-----	Listening	27/Feb/2007 11:17:52	1173	0	1173
IEXPLORE.EXE	UDP	localhost:1697	-----	Listening	27/Feb/2007 14:09:48	15	0	15

Learning Paradigm: Reinforcement Learning


- Training examples as **input**-output pairs, with **evaluative output** only



- try to increase the **reinforcement** it receives
- Example:
 - graded examinations with only overall scores but no correct answers

Example

Game playing



Black to move

Pole balancing

